Linux Academy

Hands-On Training

Configuring an Auto Scaling Application

Contents

Introduction....1
Scenario....1
Creating a Launch Configuration....1
Create an Auto Scaling Group....2
Load Balancers....5
Self-Healing Applications....5

Introduction

In this lab, we will be using concepts learned in the **AWS Certified Solutions Architect** course, in particular, the videos *Concept to Production: Building a Highly Available WordPress Application* and *Auto Scaling and Bootstrapping*.

To begin, log into the AWS Console with the credentials provided on the lab page.

Scenario

We are creating a highly-available web application called Grumpy Cat that will use AWS Auto Scaling to scale both up and down as needed, based on peak CPU loads.

Creating a Launch Configuration

Navigate to the **EC2 dashboard** from the AWS Console and select **Launch Configurations**, located in the left bar under **Auto Scaling**.

Press **Create Auto Scaling group**. AWS will provide you with a page giving you an overview of Auto Scaling group creation. Click **Create launch configuration**.

From here, you will need to select an Amazon Machine Image (AMI). On the left menu, select **My AMIs**, and check the **Shared with me** box under **Ownership**. An AMI titled *autoscaling-grumpycat* will be available from the list. Select this AMI.

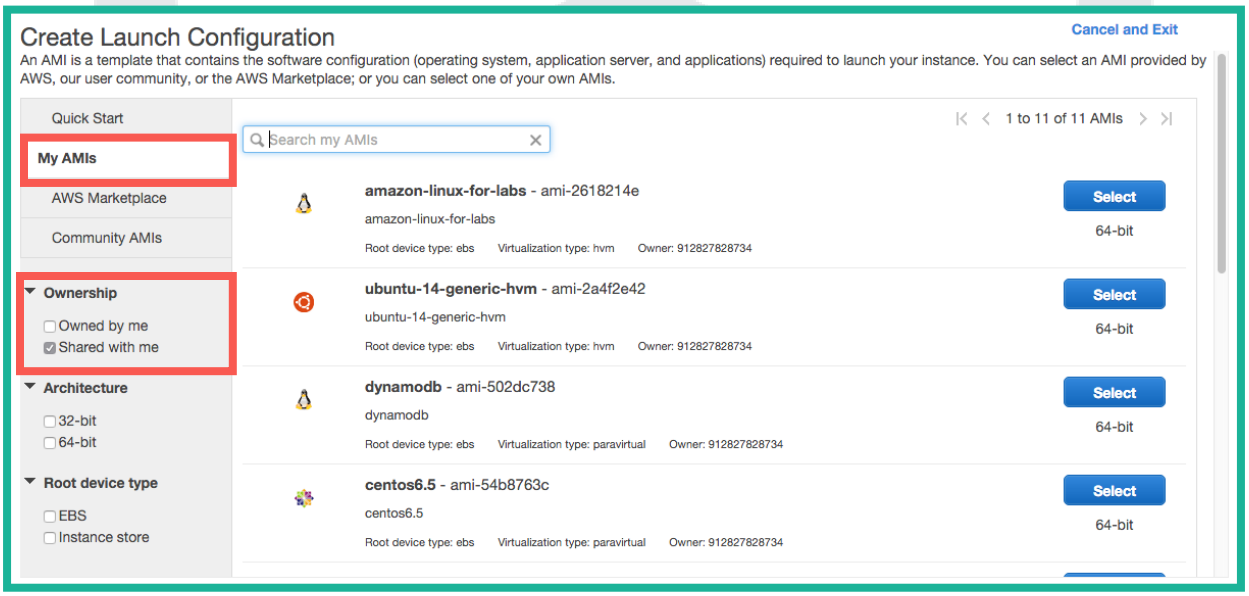

Use the default *t2.micro* plan, and press **Next: Configure details**.

We will be calling our launch configuration *GrumpyCat* (**Name**) since this is the application it will be used for. Do **not** check **Request Spot Instances**, and leave the **IAM role** set to *none*. Also, leave **Monitoring**

unchecked.

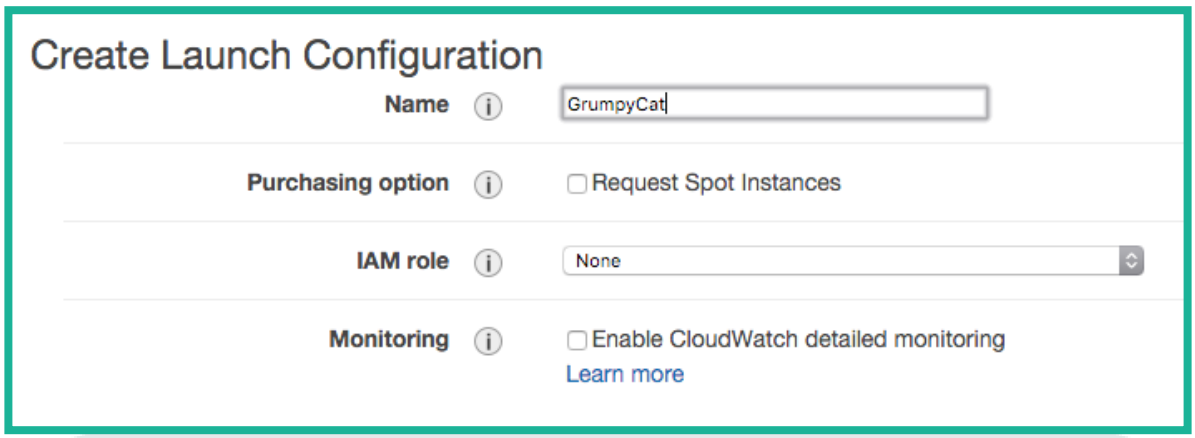

Press **Next: Add Storage**, and leave everything with the default settings. Go to **Next: Configure Security Group**.

We want to **Create a new security group**. We used *GrumpyCat* as both the security group name, and the description. If there is an option for a VPC with LinuxAcademy in the title, use this; otherwise, use the default option.

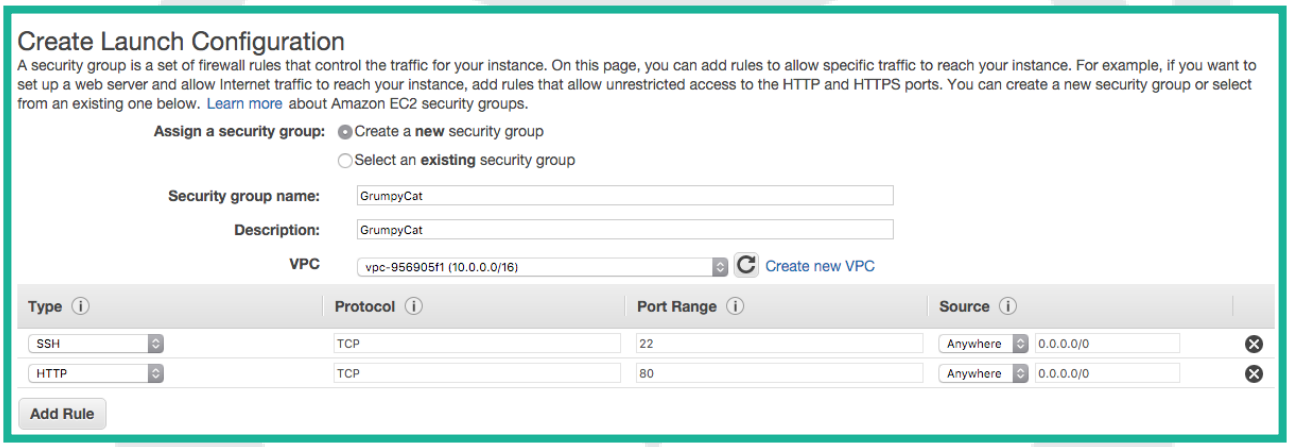

Press **Add Rule**, and add a rule for *HTTP (port 80)*. Select **Review**, and look through your settings to ensure everything is configured correctly before pressing **Create launch configuration**.

You will be prompted to select an existing or create a new key pair. Since we do not need a key pair, choose *Proceed without a key pair* from the dropdown menu, and read and check off the acknowledgment below. Press **Create launch configuration**.

Create an Auto Scaling Group

As before, we will be calling our group *GrumpyCat*. We want to start with a **Group size** of *2* instances, for high-availability purposes. Select the available VPC under **Network**, and add both subnets to the subnet area.

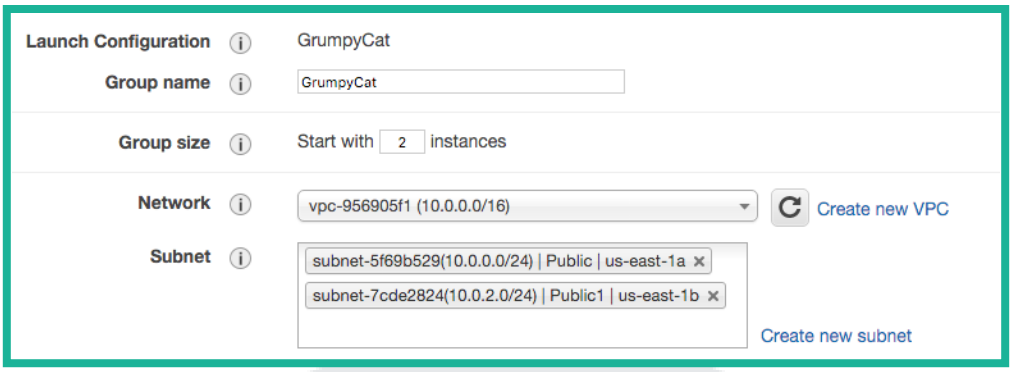

Below the subnet area is an **Advanced Details** section. Expand this so we can configure the load balancing portion of the application. Check **Receive traffic from Elastic Load Balancer(s)** and in the box below, select the single ELB available. Set the **Health Check Type** to *ELB*. You can leave the **Health Check Grace Period** at the default *300* seconds.

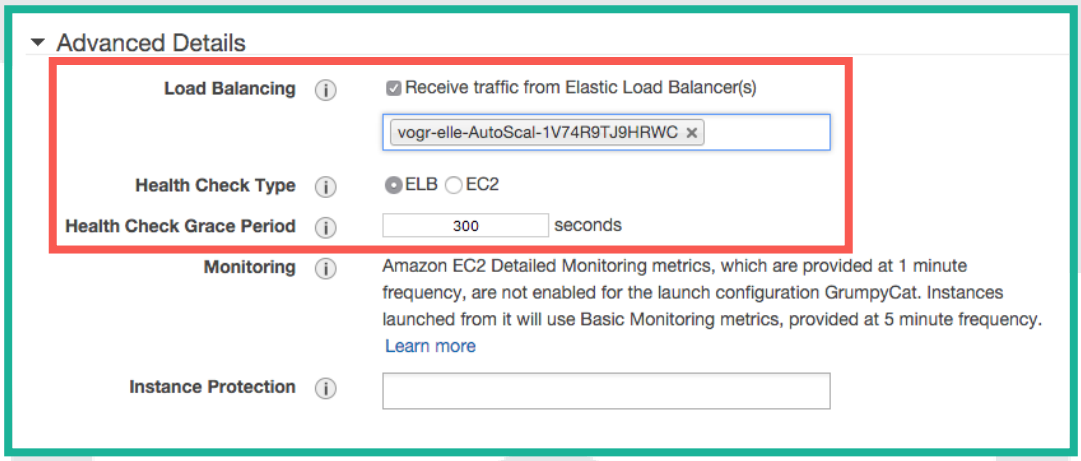

Press **Next: Configure scaling properties**. We want to **Use scaling policies to adjust the capacity of this group**. You will be presented with two options for actions and alerts: Increasing and Decreasing the group size. First, we must define the minimum and maximum amount of instances, however. Set it to **Scale between *2* and *4* instances**.

Within the **Increase Group Size** area, press **Add new alarm**. Uncheck the option to send out a notification, and change **Whenever** to be a *Maximum of CPU Utilization [that] is >= 20 Percent*. Set F**or at least** to be *1* **consecutive period(s) of** *1 Minute*. Press **Create Alarm**.

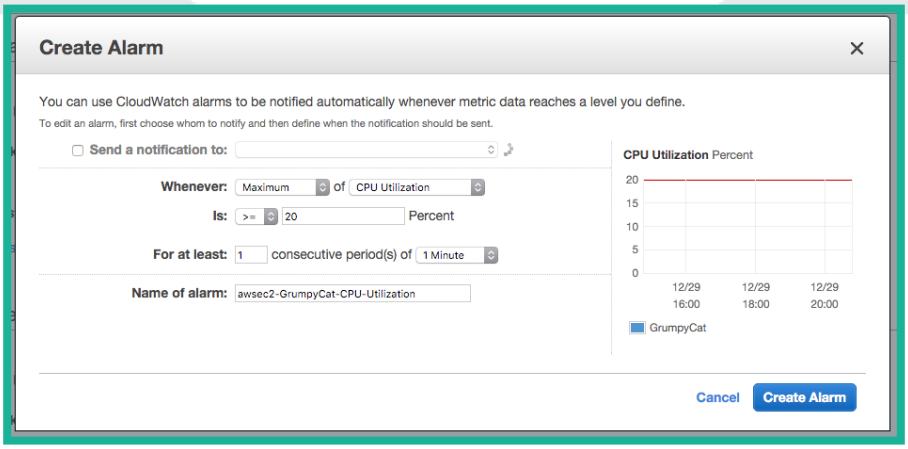

From here, we now need to define the action we want AWS to take when the alarm threshold is hit. In the **Take the action** area, we want to *Add 1* **instance**. Set the **Instances needed** to *300* **seconds to warm up after each step**.

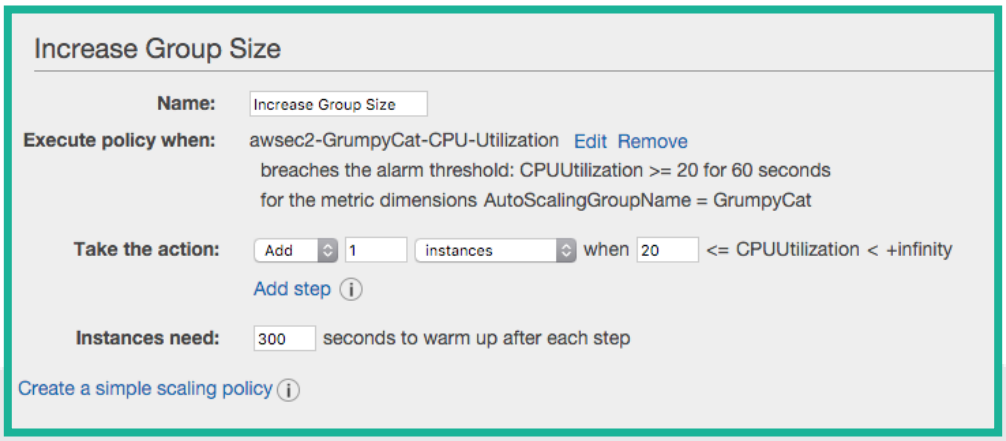

Under **Decrease Group Size**, also **Add new alarm**. Again, deselect the send notification option. Set **Whenever** to a *Minimum of CPU Utilization [that] is <= 19 Percent* **for at least** *1 consecutive period of 1 Minute*. **Create Alarm**, then set **Take the action** to *Remove 1 instances*.

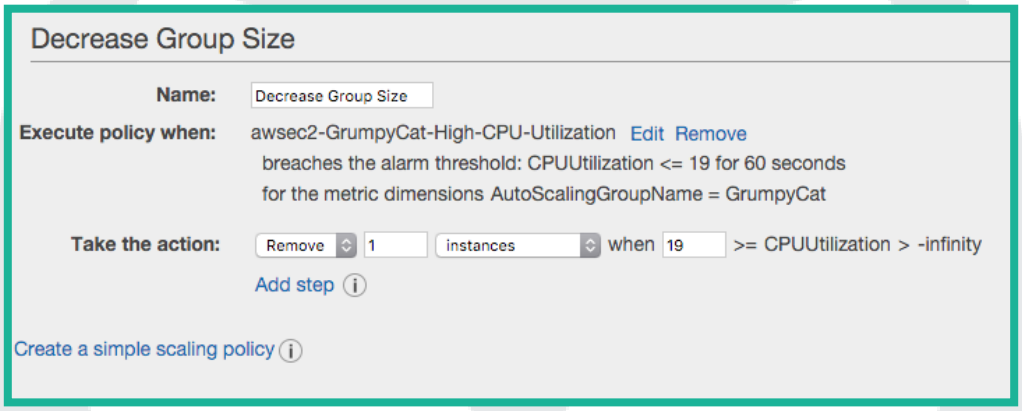

Press **Next: Configure Notifications**, ignore the warning, then press **Next: Configure Tags**. We do not need to configure tags for this instance, but to foster good habits, you may wish to give the instance a name value. Press **Review**, and review your details. Click **Create Auto Scaling group**.

If you go to your **Instances**, you will see two EC2 instances starting up.

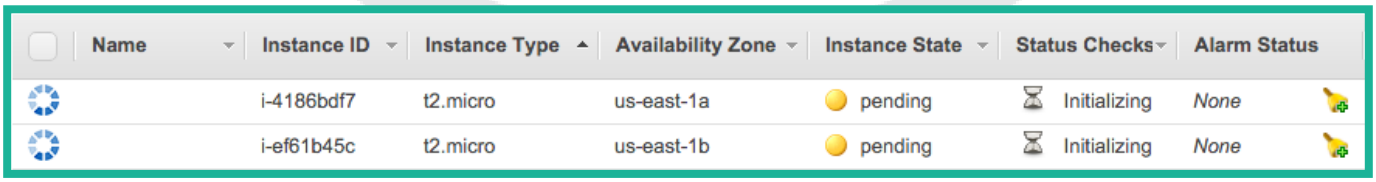

Load Balancers

If you navigate to your **Load Balancer** dashboard, located within the EC2 dashboard, under Load Balancing on the left menu, you can see the created load balancer. Select it, and then press the **Instances** tab below. Here you will see two instances, one for each zone. Wait for the **Status** of both load balancers to read *InService* before continuing. This can take up to ten minutes.

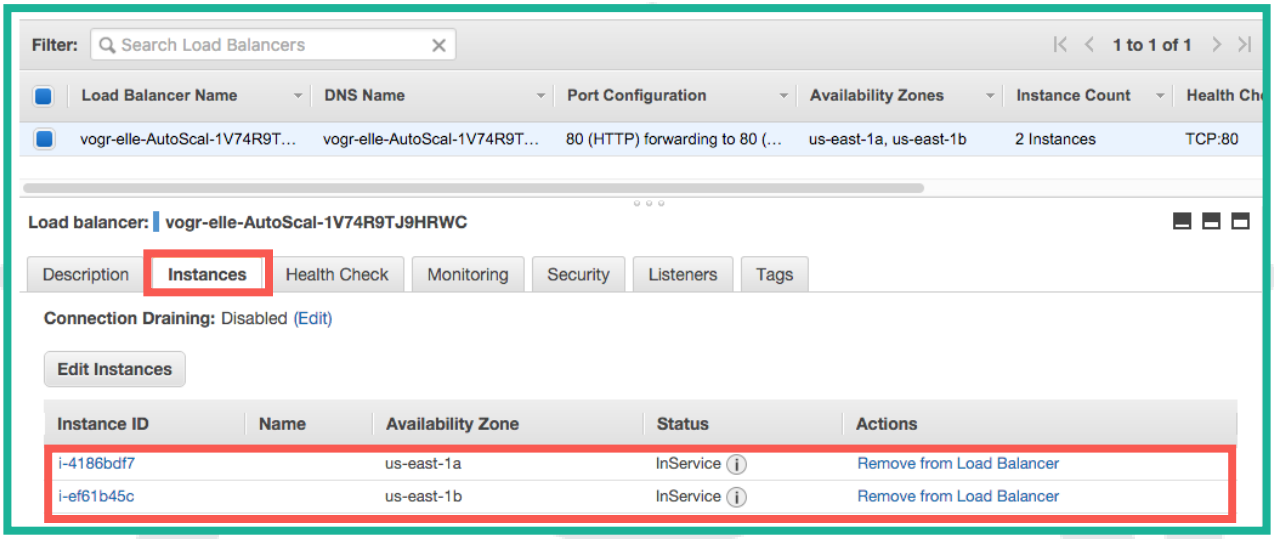

Return to the **Description** tab, and copy and **DNS Name**, then paste it into your web browser. You should see your Grumpy Cat.

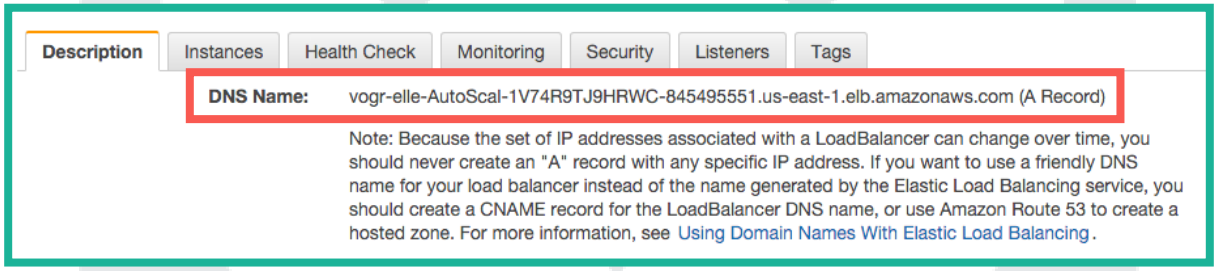

Self-Healing Applications

To test if our instance is truly self-healing, we will now need to return to **Instances**, under the left menu. Here, we need to terminate one of the running instances: If our application is self-healing, another instance should be brought up, once the initial instance is terminated. You may also terminate the other running instance if so desired. Because we created rules that call for a minimum of two instances, that, too, will bring up a replacement instance. You may need to refresh the page to witness this.

Return to the link in your web browser and refresh to ensure Grumpy Cat is still available, even after your instances have been deleted and recreated.